

In [200...

```
#!pip install -U scikit-learn
#!pip install numpy
#!pip install pandas
#!pip install matplotlib
#!pip install seaborn
```

In [200...

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import math
```

In [200...

```
df=pd.read_csv(r"C:\Users\vaibhav kumar\Downloads\Boston_Train_Test.csv", encoding= 'unicode_escape')
df
```

Out[200...

	Unnamed: 0	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat	medv
0	0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
1	1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
2	2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
3	3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
4	4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2
...
501	501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	21.0	391.99	9.67	22.4
502	502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	21.0	396.90	9.08	20.6
503	503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	21.0	396.90	5.64	23.9
504	504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	21.0	393.45	6.48	22.0
505	505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	21.0	396.90	7.88	11.9

506 rows × 15 columns

In [201...

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Unnamed: 0   506 non-null    int64
1   crim        506 non-null    float64
2   zn          506 non-null    float64
3   indus       506 non-null    float64
4   chas        506 non-null    int64
5   nox         506 non-null    float64
6   rm          506 non-null    float64
7   age         506 non-null    float64
8   dis         506 non-null    float64
9   rad         506 non-null    int64
10  tax         506 non-null    int64
11  ptratio     506 non-null    float64
12  black       506 non-null    float64
13  lstat       506 non-null    float64
```

14 medv 506 non-null float64
dtypes: float64(11), int64(4)
memory usage: 59.4 KB

In [201...

```
df.describe()
```

Out[201...

	Unnamed: 0	crim	zn	indus	chas	nox	rm	age	dis	
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	252.500000	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	3.795043	
std	146.213884	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2.105710	
min	0.000000	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.129600	
25%	126.250000	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2.100175	
50%	252.500000	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	3.207450	
75%	378.750000	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5.188425	
max	505.000000	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.126500	

In [201...


```
df.corr()
```

Out[201...

	Unnamed: 0	crim	zn	indus	chas	nox	rm	age	dis	rad
Unnamed: 0	1.000000	0.407407	-0.103393	0.399439	-0.003759	0.398736	-0.079971	0.203784	-0.302211	0.686002
crim	0.407407	1.000000	-0.200469	0.406583	-0.055892	0.420972	-0.219247	0.352734	-0.379670	0.625505
zn	-0.103393	-0.200469	1.000000	-0.533828	-0.042697	-0.516604	0.311991	-0.569537	0.664408	-0.311948
indus	0.399439	0.406583	-0.533828	1.000000	0.062938	0.763651	-0.391676	0.644779	-0.708027	0.595129
chas	-0.003759	-0.055892	-0.042697	0.062938	1.000000	0.091203	0.091251	0.086518	-0.099176	-0.007368
nox	0.398736	0.420972	-0.516604	0.763651	0.091203	1.000000	-0.302188	0.731470	-0.769230	0.611441
rm	-0.079971	-0.219247	0.311991	-0.391676	0.091251	-0.302188	1.000000	-0.240265	0.205246	-0.209847
age	0.203784	0.352734	-0.569537	0.644779	0.086518	0.731470	-0.240265	1.000000	-0.747881	0.456022
dis	-0.302211	-0.379670	0.664408	-0.708027	-0.099176	-0.769230	0.205246	-0.747881	1.000000	-0.494588
rad	0.686002	0.625505	-0.311948	0.595129	-0.007368	0.611441	-0.209847	0.456022	-0.494588	1.000000
tax	0.666626	0.582764	-0.314563	0.720760	-0.035587	0.668023	-0.292048	0.506456	-0.534432	0.910228
ptratio	0.291074	0.289946	-0.391679	0.383248	-0.121515	0.188933	-0.355501	0.261515	-0.232471	0.464741
black	-0.295041	-0.385064	0.175520	-0.356977	0.048788	-0.380051	0.128069	-0.273534	0.291512	-0.444413
lstat	0.258465	0.455621	-0.412995	0.603800	-0.053929	0.590879	-0.613808	0.602339	-0.496996	0.488676
medv	-0.226604	-0.388305	0.360445	-0.483725	0.175260	-0.427321	0.695360	-0.376955	0.249929	-0.381626

In [201...

```
sns.heatmap(df.corr())  
plt.show()
```



X y Split

```
In [201... X = df.drop("medv", axis = 1)
y = df["medv"]
```

Test Train Splitting

```
In [201... from sklearn.model_selection import train_test_split
```

```
In [201... X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

```
In [201... X_train
```

```
Out[201...
   Unnamed: 0  crim  zn  indus  chas  nox  rm  age  dis  rad  tax  ptratio  black  lstat
240          240  0.11329  30.0   4.93    0  0.428  6.897  54.3  6.3361  6  300    16.6  391.25  11.38
380          380  88.97620   0.0  18.10    0  0.671  6.968  91.9  1.4165  24  666    20.2  396.90  17.21
212          212  0.21719   0.0  10.59    1  0.489  5.807  53.8  3.6526  4  277    18.6  390.94  16.03
  2           2  0.02729   0.0   7.07    0  0.469  7.185  61.1  4.9671  2  242    17.8  392.83   4.03
104          104  0.13960   0.0   8.56    0  0.520  6.167  90.0  2.4210  5  384    20.9  392.69  12.33
...          ...    ...    ...    ...    ...    ...    ...    ...    ...    ...    ...    ...    ...
458          458  7.75223   0.0  18.10    0  0.713  6.301  83.7  2.7831  24  666    20.2  272.21  16.23
 91           91  0.03932   0.0   3.41    0  0.489  6.405  73.9  3.0921  2  270    17.8  393.55   8.20
390          390  6.96215   0.0  18.10    0  0.700  5.713  97.0  1.9265  24  666    20.2  394.43  17.11
295          295  0.12932   0.0  13.92    0  0.437  6.678  31.1  5.9604  4  289    16.0  396.90   6.27
 23           23  0.98843   0.0   8.14    0  0.538  5.813  100.0  4.0952  4  307    21.0  394.54  19.88
```

354 rows × 14 columns

```
In [201... y_train
```

```
Out[201...] 240    22.0
             380    10.4
             212    22.4
              2    34.7
             104    20.1
             ...
            458    14.9
             91    22.0
            390    15.1
            295    28.6
             23    14.5
Name: medv, Length: 354, dtype: float64
```

```
In [201...] # Splliting of Train data is done randomly to avoid potential selection biases arising. This is be
            # non-biased resuts, it also means that results can differ from run to run.
```

```
In [202...] X_test
```

```
Out[202...] Unnamed: 0    crim    zn    indus    chas    nox    rm    age    dis    rad    tax    ptratio    black    lstat
504          504  0.10959    0.0    11.93     0  0.573  6.794  89.3  2.3889     1   273      21.0  393.45    6.48
393          393  8.64476    0.0    18.10     0  0.693  6.193  92.6  1.7912    24   666      20.2  396.90   15.17
248          248  0.16439   22.0     5.86     0  0.431  6.433  49.1  7.8265     7   330      19.1  374.71    9.52
134          134  0.97617    0.0    21.89     0  0.624  5.757  98.4  2.3460     4   437      21.2  262.76   17.31
317          317  0.24522    0.0     9.90     0  0.544  5.782  71.7  4.0317     4   304      18.4  396.90   15.94
...          ...    ...    ...    ...    ...    ...    ...    ...    ...    ...    ...    ...    ...    ...
343          343  0.02543   55.0     3.78     0  0.484  6.696  56.4  5.7321     5   370      17.6  396.90    7.18
61           61  0.17171   25.0     5.13     0  0.453  5.966  93.4  6.8185     8   284      19.7  378.08   14.44
105          105  0.13262    0.0     8.56     0  0.520  5.851  96.7  2.1069     5   384      20.9  394.05   16.47
499          499  0.17783    0.0     9.69     0  0.585  5.569  73.5  2.3999     6   391      19.2  395.77   15.10
450          450  6.71772    0.0    18.10     0  0.713  6.749  92.6  2.3236    24   666      20.2    0.32   17.44
```

152 rows × 14 columns

```
In [202...] y_test
```

```
Out[202...] 504    22.0
             393    13.8
             248    24.5
             134    15.6
             317    19.8
             ...
            343    23.9
             61    16.0
            105    19.5
            499    17.5
            450    13.4
Name: medv, Length: 152, dtype: float64
```

```
In [202...] # Splliting of Test data is done randomly to avoid potential selection biases arising. This is ber
            # non-biased resuts, it also means that results can differ from run to run.
```

Normalization

```
In [202... from sklearn.preprocessing import MinMaxScaler  
sc = MinMaxScaler()
```

```
In [202... sc.fit(X_train)  
X_train = sc.transform(X_train)  
X_test = sc.transform(X_test)
```

```
In [202... X_train.shape
```

```
Out[202... (354, 14)
```

```
In [202... X_train
```

```
Out[202... array([[0.47524752, 0.00120232, 0.3      , ..., 0.42553191, 0.98567372,  
        0.27375887],  
       [0.75247525, 1.          , 0.          , ..., 0.80851064, 1.          ,  
        0.43914894],  
       [0.41980198, 0.00237013, 0.          , ..., 0.63829787, 0.98488767,  
        0.40567376],  
       ...,  
       [0.77227723, 0.07818185, 0.          , ..., 0.80851064, 0.993737  ,  
        0.43631206],  
       [0.58415842, 0.00138249, 0.          , ..., 0.36170213, 1.          ,  
        0.12879433],  
       [0.04554455, 0.01103868, 0.          , ..., 0.89361702, 0.99401592,  
        0.51489362]])
```

```
In [202... X_test
```

```
Out[202... array([[ 0.9980198 , 0.00116073, 0.          , ..., 0.89361702,  
        0.99125209, 0.13475177],  
       [ 0.77821782, 0.09709398, 0.          , ..., 0.80851064,  
        1.          , 0.3812766  ],  
       [ 0.49108911, 0.00177667, 0.22       , ..., 0.69148936,  
        0.94373447, 0.22099291],  
       ...,  
       [ 0.20792079, 0.00141958, 0.          , ..., 0.88297872,  
        0.99277347, 0.41815603],  
       [ 0.98811881, 0.00192773, 0.          , ..., 0.70212766,  
        0.99713474, 0.37929078],  
       [ 0.89108911, 0.07543452, 0.          , ..., 0.80851064,  
       -0.00557838, 0.44567376]])
```

```
In [202... X_test.shape
```

```
Out[202... (152, 14)
```

Linear Regression

```
In [202... from sklearn.linear_model import LinearRegression  
reg = LinearRegression()
```

```
In [203... x = df["nox"]
```

```
y = df["dis"]
```

```
In [203... from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

```
In [203... X_train = np.array(X_train).reshape(-1, 1)  
X_test = np.array(X_test).reshape(-1, 1)
```

```
In [203... #X_train
```

```
In [203... model = reg.fit(X_train, y_train)
```

```
In [203... model.coef_
```


```
Out[203... array([-13.58585091])
```

```
In [203... model.intercept_
```

```
Out[203... 11.268763847893577
```

```
In [203... y_pred = model.predict(X_test)
```

```
In [203... plt.scatter(X_test, -y_pred, c = "r")  
plt.scatter(X_test, -y_test, c = "b")  
plt.show()
```



```
In [203... from sklearn.metrics import mean_absolute_error  
from sklearn.metrics import mean_squared_error  
from sklearn.metrics import r2_score
```

```
In [204... mean_absolute_error(y_test, y_pred)
```

```
Out[204... 1.1609318877854597
```

```
In [204... mean_squared_error(y_test, y_pred)
```

Out[204...] 2.281290171525275

In [204...] `np.sqrt(mean_squared_error(y_test, y_pred))`

Out[204...] 1.510394045117126

In [204...] `p=r2_score(y_test, y_pred)`

In [204...] `p`

Out[204...] 0.5644595318232439

Multiple Linear Regression

In [204...] `x = df.drop("medv", axis = 1)`
`y = df["medv"]`

In [204...] `x`

Out[204...]

	Unnamed: 0	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat
0	0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98
1	1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14
2	2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03
3	3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94
4	4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33
...
501	501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	21.0	391.99	9.67
502	502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	21.0	396.90	9.08
503	503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	21.0	396.90	5.64
504	504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	21.0	393.45	6.48
505	505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	21.0	396.90	7.88

506 rows × 14 columns

In [204...] `from sklearn.model_selection import train_test_split`
`X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=55)`

In [204...] `reg = LinearRegression()`
`model = reg.fit(X_train, y_train)`

In [204...] `model.coef_`

Out[204...] `array([-7.60758272e-04, -1.33440243e-01, 5.76329038e-02, 3.54041723e-02,`
 `1.69281625e+00, -1.86212072e+01, 3.56460749e+00, -1.74212324e-03,`

```
-1.58644494e+00, 3.26202485e-01, -1.27117387e-02, -9.64352356e-01,  
8.58633637e-03, -5.37874887e-01])
```

```
In [205... model.intercept_
```

```
Out[205... 39.62322051209671
```

```
In [205... y_pred = model.predict(X_test)
```

```
In [205... mean_absolute_error(y_test, y_pred)
```

```
Out[205... 2.9495942985662755
```

```
In [205... mean_squared_error(y_test, y_pred)
```

```
Out[205... 17.836167010161777
```

```
In [205... np.sqrt(mean_squared_error(y_test, y_pred))
```

```
Out[205... 4.223288648690944
```

```
In [205... r2_score(y_test, y_pred)
```

```
Out[205... 0.753625070837884
```

Polynomial Regression

```
In [205... X = df["nox"]  
y = df["dis"]
```

```
In [205... X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

```
In [205... p_test = X_test
```

```
In [205... from sklearn.preprocessing import PolynomialFeatures
```

```
In [206... poly = PolynomialFeatures(2)
```

```
In [206... X_train = np.array(X_train).reshape(-1, 1)  
X_test = np.array(X_test).reshape(-1, 1)
```

```
In [206... poly.fit(X_train)
```

```
Out[206... ▼ PolynomialFeatures  
PolynomialFeatures()
```



```
In [206... X_train = poly.transform(X_train)
X_test = poly.transform(X_test)
```

```
In [206... #X_test
```

```
In [206... reg = LinearRegression()
```

```
In [206... model = reg.fit(X_train, y_train)
```

```
In [206... model.coef_
```


```
Out[206... array([ 0.          , -73.89605897,  50.02247161])
```

```
In [206... model.intercept_
```

```
Out[206... 28.746142405931856
```

```
In [206... y_pred = model.predict(X_test)
```

```
In [207... plt.scatter(X_test[:, 1], y_pred, c = "r")
plt.scatter(X_test[:, 1], y_test, c = "b")
plt.show()
```



```
In [207... mean_absolute_error(y_test, y_pred)
```

```
Out[207... 0.8763119152067171
```

```
In [207... mean_squared_error(y_test, y_pred)
```

```
Out[207... 1.3003165078669456
```

```
In [207... np.sqrt(mean_squared_error(y_test, y_pred))
```

```
Out[207... 1.140314214533409
```

```
In [207...
```

```
r2_score(y_test, y_pred)
```

Out[207...] 0.7193327117492254

Observations

In [207...] *# In case of Multiple Liner Regression the r2 value comes near to 0.754 that
means 75% data is reliable so house prediction(Medv) in Boston.
It is nearly a Good data but not perfect.*

THANK YOU